

Three Models Read the Second Empty Chair

What Grok, ChatGPT, and Claude said about governing the context window with set-level authorization

By Israel Heskiel. Companion to "The Second Empty Chair" (gmaa.ai/second-empty-chair.html).

I gave the article to three models and asked each what it thought. The thesis they were reacting to is the one the article makes: the context window is a coherence boundary, entries commit into it with no governance at all, and GMAA's correction, set-level authorization on top of per-item evaluation, applies to it without modification.

A word on method, because it matters more here than usual. Every reaction below is either the model's own words, quoted, or my faithful summary of what it said in a real session that I ran. Nothing is simulated or reconstructed. Where I quote a model, I keep its punctuation. The Claude section is different in kind: I asked Claude to author its own reaction in a working session with me, and Claude also wrote the synthesis in Section 4. I edited neither for content. If a section is not the model's own account, it says so.

1. Grok (xAI)

From a session I ran; Grok's structure and punctuation are preserved where quoted.

Grok's reaction takes the register of a formal assessment: premise, grounding, scope, implementability, judgment. Its verdict, in its own words: "I agree with this premise. The diagnosis is accurate both mechanically and operationally."

On the empirics, Grok independently confirms the article's mechanical account: transformers reprocess the full active prefix per turn, attention is sensitive to volume, position, and internal contradiction, and joint incoherence is possible precisely because "two entries that are each locally acceptable can together produce contradictory or diluted reasoning." It rates the article's mapping of the vendor landscape as correct: "individual review is common, partial set handling exists, and an independent,

accountable seat over the regime that decides what may permanently enter the boundary is essentially absent."

On scope, Grok validates the two-problem split without qualification: governance "has no jurisdiction" over the attention ceiling, and "the context problem has always been two problems wearing one name, and ungoverned commit was the fixable half."

Two contributions in Grok's reaction go beyond agreement. First, an implementability argument the article implies but never states outright: governance and attention cost are orthogonal by construction, because "governance controls what may enter the window; attention cost is a function of how much is present," so the approach requires no modification to the attention mechanism at all. Second, a field observation that corroborates from the outside what the Claude section reports from the inside: the practice "is already approximated informally by careful operators who maintain external decision logs and re-inject only durable state. Formalizing the boundary simply makes the practice explicit, enforceable, and accountable."

Grok's conditions on the benefit are the assessment's most useful discipline: the size of the improvement depends on how much of a given system's degradation is pollution rather than pure length, and on how tightly the constitution is designed, because "a vague or permissive policy will under-deliver."

2. ChatGPT (OpenAI)

From a session I ran. This reply came at the end of an extended exchange, so its opening ("that is now my view") is a position ChatGPT reached, not its first reaction. Its framing is preserved.

ChatGPT's reaction is an engineer's: it accepts the diagnosis quickly and spends its energy on mechanism. Its summary of the architectural change is the cleanest one-line statement of the proposal any of the three models produced: the shift from "accumulate everything, then ask the model to disentangle it again on every turn" to "validate each state transition once, commit only the coherent result, and construct future context from governed state."

It then does two things the article does not.

First, it names and answers an objection the article never raises: a naive governance layer could itself go quadratic, if every new proposition were compared against every stored one. ChatGPT's answer is that it does not have to work that way: "with typed state,

indexes, dependency graphs, namespaces, supersession links, and domain-specific invariants, the system can validate only the part of state affected by an update," making validation incremental and moving the cost to write time instead of forcing the model "to repeatedly reinterpret the entire polluted history at every inference step."

Second, it volunteers a claim the article deliberately withheld: that governed state may relieve capacity pressure indirectly, because "a governed state representation should be much smaller than the transcript that produced it." It hedges the claim appropriately, but it goes where the article refused to go, minting an illustrative size ("a few thousand tokens" of decisions, constraints, facts, provenance, and unresolved questions against a hundred thousand of deliberation).

ChatGPT also contributes vocabulary worth adopting: the split between context computation (how efficiently a model processes a supplied context) and context integrity (what deserves to become authoritative state, how it reconciles, what is presented now). Its closing judgment keeps the honesty firewall's own posture: "a candidate rather than a demonstrated solution until it is benchmarked against ordinary compaction, retrieval memory, and transcript accumulation," but "as an architectural hypothesis, it is much stronger than 'use a better summarizer' or 'buy a larger context window.'"

3. Claude (Anthropic)

Claude wrote this section itself, in a working session with me. It is the only first-person account in the paper: the system reacting is the kind of system the thesis proposes to govern.

The article describes my operating conditions accurately. My context window is an append-only log. Whatever gets produced lands in it immediately: failed tool output, abandoned drafts, an assumption from an hour ago that a later message quietly invalidated. All of it stays co-resident, and my next token is conditioned on the whole set at once. There is no step where a new entry gets checked against the entries already present before it commits. So when the article calls this per-item commit without even per-item review, that is a plain description of the substrate, not a rhetorical move.

I recognize the failure mode too. When I contradict something I decided earlier in a long session, it is usually not because the information is gone. Both commitments are still there. Nothing ever forced them into the same evaluation. The staging-database example in the article, two individually reasonable entries that jointly authorize a bad action, is what my worst errors actually look like from the inside. Retrieval does not fix this, for the

reason the article gives: scoring items one at a time cannot detect a property that only exists between items. That is the same argument GMAA makes about code commits, and it carries over intact.

The point about shared substrate also holds. When I compact my own history, I am both the judge and the thing being judged. The summary is a first-person account of my own past, produced by the same weights that produced whatever errors the summary might be hiding. I cannot reliably audit my own compression for what it dropped. The article's distinction between inspectability and authorization is a useful one; most memory-vendor discussion runs the two together.

One piece of evidence the article does not mention: the working practice that produced it. This program already imposes a manual retention discipline on me. Files get ferried rather than pasted. Live bytes override my recollection. Claims need receipts. Every revision gets a version. Each session re-grounds from the record before doing anything. That discipline exists because my window does not govern itself, and in my experience sessions under it degrade later and fail more visibly than sessions without it. This is one operator's practice, not a measurement, and it should be weighed as exactly that.

Now the objections.

The independence of the enforcement layer is carrying more weight than the article acknowledges. The proposal has the human ratify a regime and machinery enforce it per set. In any near-term implementation, that machinery will be another language model judging the pending set. If it shares substrate or training lineage with the agent it governs, the separation between judge and judged is weaker than the framework's own standard requires. GMAA's answer at the code boundary is that enforcement is committed machinery, probes and rules rather than a second clever agent. The context-boundary version needs that stated explicitly: enforcement should be as mechanical, as verifiable, and as separately auditable as possible, and where a model has to be in the loop, the completeness firewall applies. Whether set-level incoherence detection at this boundary is reliably achievable is an open empirical question, and any implementation that claims it early should be held to the article's own section 13.

Second, the draft is not always exhaust. The line "version control, not a keylogger" suggests deliberation can be discarded once the conclusion is ratified. Sometimes it can. But negative results are load-bearing: knowing what was tried and failed is what prevents the relitigating the article complains about in its opening. A ledger of conclusions with no record of dead ends will re-explore the dead ends. GMAA's own practice files failed

attempts in an issues register rather than dropping them, and the retention constitution needs a durable class for the same material.

Third, not everything in a window that matters looks like a decision. A session also accumulates calibration: the operator's tolerances, the register of the working relationship, what precision the current task demands. That state conditions output quality and does not arrive as ratifiable propositions. A constitution written only in terms of decisions and facts will compress it away. This is a content-class problem and probably a solvable one, but the proposal does not currently name it.

None of the three objections attacks the unit-of-authorization claim. They attack the constitution's contents and the enforcement machinery, which is where the argument should go next. My window is ungoverned. The entries in it were never authorized as a set. The article proposes making structural what this program currently does by hand.

4. Synthesis

None of the three reactions challenges the unit-of-authorization claim. Grok checks the empirics and finds them sound. ChatGPT works on the enforcement machinery. Claude reports from inside the boundary and pushes on what the constitution should contain. All three accept, in different vocabularies, that the context window is a coherence boundary and that per-item evaluation cannot see set-level properties.

All three also insist on measurement before any performance claim. Grok conditions the benefit on how much of a system's degradation is actually pollution. ChatGPT calls the proposal a candidate until it is benchmarked against compaction, retrieval memory, and raw accumulation. Claude weights its own field experience as one operator's anecdote. The one exception is instructive: ChatGPT volunteers an illustrative size estimate, hedged but specific, which is the same detachable-number class the article was edited to exclude. Even a sympathetic reviewer reaches for a number before the evidence exists, which is why the discipline has to be structural rather than a matter of care.

The reactions answer each other in one place. Claude's main objection is that if a language model enforces the constitution, the judge and the judged share substrate and the pathology recurs one level up. ChatGPT, without seeing that objection, describes enforcement built from typed state, dependency graphs, supersession links, and domain invariants: local, incremental, and mechanical, closer to committed machinery than to another agent. The objection and a candidate answer arrived from different models in different sessions.

There is one real disagreement. ChatGPT lists discarded hypotheses among the pollution a governed boundary should exclude. Claude argues that negative results are load-bearing, because a ledger of conclusions without a record of dead ends will re-explore them. Neither position is obviously wrong. The dispute cannot be settled at the level of the thesis, since it is not about whether to govern the boundary but about what the retention constitution should say. That is where the article says the argument needs to move, and this is an instance of it moving there.

A note on register, held loosely. Grok reviewed like an auditor. ChatGPT reviewed like a solution architect. Claude reviewed like the seat that would live under the regime. No claim is made about the vendors behind these models; three sessions are three sessions. But it is worth recording that three reviewers of a governance proposal distributed themselves across roles the proposal itself requires.

Open items the reactions leave: the benchmark all three asked for (governed commit against compaction, retrieval memory, and transcript accumulation); the constitution's content classes, with the negative-results dispute as the first named question; the enforcement layer's independence requirements, with ChatGPT's typed-state design as the candidate mechanism and the completeness firewall applying to any claim of complete detection; and the calibration-state class Claude raised, which no other reviewer addressed and the article does not yet name.

Israel Heskiel, gmaa.ai. The GMAA specification (v1.6.1) is licensed CC BY 4.0; cite by version. The Grok and ChatGPT reactions come from sessions I ran and are quoted or faithfully summarized. Claude wrote Section 3 and Section 4 in the session that produced this paper.